

Grounded Extraction of Technology Mentions from Papers and Patents

Abstract

Tracking how technologies emerge and spread begins with reading individual documents and naming the technologies in them. We argue that this apparently routine extraction task has an unusual error profile once its output is aggregated into trends: a fabricated technology is not averaged away by scale, because a spurious term whose rate follows corpus growth produces a curve indistinguishable from a genuine emerging technology, while a missed mention mostly rescales a series without changing its shape. We therefore build an extractor around structural constraints rather than around a confidence threshold. Every emitted mention is a pair of character offsets into its source document with a cited evidence sentence, so a model that names a technology we cannot locate has named nothing; candidates are generated by five deliberately over-eager recallers and filtered by a guard stack, an independent verifier cascade and a single calibrated abstention threshold. On a 895-mention gold set spanning 24 arXiv abstracts and 20 USPTO publications across nine technology areas, the full system reaches 0.548 partial F_1 at a measured hallucination rate of 0.000. An ablation over six components finds only one of them unambiguously load-bearing — the verifier — while two others buy strict precision at a measurable cost in overlap recall and in ranking quality, which we report rather than smooth over. We also report a negative result: span grounding bounds but does not close indirect prompt injection, because an adversary can write their phantom technology into their own patent.

1 Introduction

A lab that wants to track the evolution of technology through the scientific and patent literature needs, before anything else, to know which technologies each document is about. The rest of the pipeline — time series, emergence detection, successor analysis — consumes that output and cannot

repair it.

The task looks like scientific named-entity recognition, and the obvious move is to apply the machinery built for that: a span tagger over a scientific ontology (Augenstein et al., 2017; Luan et al., 2018), a domain-pretrained encoder (Beltagy et al., 2019), a span-enumeration model (Wadden et al., 2019; Zhong and Chen, 2021), or a prompted generative extractor. Each of these is a reasonable starting point and none of them is designed for the property that matters here.

That property is an asymmetry between the two kinds of error. Suppose an extractor invents a technology — a plausible name, a plausible type — at some low rate. Downstream, trend detection works by comparing a term’s frequency against a baseline, so a constant spurious rate produces a flat line and does no harm, but a spurious rate that tracks corpus growth produces a rising line that is indistinguishable from a real emerging technology and that nobody will go back and check. A missed mention behaves quite differently: it usually rescales a series without changing its shape. Precision errors are therefore not merely more costly than recall errors, they are qualitatively different, and they are the errors that generative extractors make (Ji et al., 2023).

We take the design consequence seriously. Rather than tune an extractor to be cautious, we make the failure unrepresentable: a mention is defined as a pair of character offsets into the source document, its surface form is the slice at those offsets, and it must cite a sentence from the document that contains it. A model that returns a technology name we cannot locate has returned nothing. This is the span-level analogue of attribution work in generation (Rashkin et al., 2023; Min et al., 2023), applied as a hard gate rather than as a post-hoc score.

Structural grounding costs recall — it discards a correct identification that was paraphrased — so

the rest of the system is built to give that recall back somewhere it can be checked. Five recallers over-generate on purpose, a guard stack removes what is provably wrong, an independent verifier cascade challenges what survives, and a single calibrated threshold decides what to withhold for human review. Contributions:

1. A precision-dominant formulation of technology extraction for evolution tracking, with selective prediction (El-Yaniv and Wiener, 2010; Geifman and El-Yaniv, 2017) as its evaluation frame rather than F_1 at full coverage.
2. A pipeline in which grounding, evidence and provenance are checked invariants, evaluated by ablation on both genres.
3. A 895-mention gold set over 44 papers and patents spanning nine technology areas, with annotation guidelines and its limitations stated.
4. An adversarial test suite that drives the pipeline with a deliberately hostile model, including the negative result on prompt injection in §6.

2 What makes this task not standard NER

The vocabulary is open by construction. The point of tracking technology is to see new technology. An extractor whose output is constrained to an ontology can never report one, which rules out entity linking as the primary mechanism and makes a knowledge base an anchor rather than a target. In our gold set, 55% of mentions are specific artefacts, and the ones that matter most for tracking — the document’s own contribution — are by definition the ones least likely to be in any lexicon.

Granularity is the whole problem in disguise. “Machine learning”, “attention mechanism” and “FlashAttention” are all technologies at different levels. A tracker that mixes them reports that machine learning is the fastest growing technology of every year since 1995. We therefore type mentions over a nine-way ontology that separates artefacts, methods and materials (tracked) from fields, tasks, metrics, datasets and tools (extracted, annotated, but excluded from the tracked stream), and record a granularity level.

Stance carries most of the signal. A mention count conflates a technology being introduced,

used and beaten. We tag each mention as PROPOSED, USED, COMPARED, BACKGROUND or — for patents — CLAIMED, because the shift of a technology’s role distribution from proposed to used is the clearest single indicator of adoption.

The two genres fail differently. Papers are fluent, citation-dense and name their contribution explicitly. Patents are deliberately abstract, repeat component names dozens of times, and bury the invention in claim language whose purpose is breadth rather than clarity. A lexicon tuned on one is close to useless on the other; §5 quantifies this.

3 Approach

3.1 Structure-aware ingestion

Normalisation repairs ligatures, line-break hyphenation and typographic characters, and masks citation markers, while maintaining a per-character map back into the raw text. Invertibility is not a nicety: when a human audits an extraction, they must be shown the mention as it appears in the original filing.

Both genres are segmented onto one closed set of structural zones — title, abstract, method, related work for papers; field, background, summary, detailed description and individual claims for patents — so that a single stance classifier serves both. Patent claims are split individually and each is classified as independent or dependent by detecting back-references (of claim 1), which matters because a term in an independent claim is being monopolised, not merely discussed.

3.2 Recall: over-generate on purpose

Five recallers run over each document, in three cost tiers.

Noun-phrase lattice. Dependency chunks (Honnibal et al., 2020) are expanded into a small span lattice: every left truncation after leading vague modifiers, crossed with alternative noun-final right boundaries. Two failures forced the second direction. Whether the term is “string similarity” or “superficial string similarity” depends on information the recaller does not have; and the parser attaches a following word often enough to matter (“obstacle-aware route planning grounds”), which no left truncation recovers.

Abbreviation mining. Schwartz–Hearst (Schwartz and Hearst, 2003) recovers the defining parenthetical mention, which is both a high-precision candidate and the best within-document

coreference signal available anywhere in the pipeline.

Terminology statistics. C-value/NC-value (Frantzi et al., 2000) is the only recaller sensitive to *nesting*, which is exactly the granularity problem; we add a termhood penalty in the spirit of Kageura and Umino (1996).

Knowledge base. Aho–Corasick matching against the Computer Science Ontology (Salatino et al., 2020). Structurally unable to find anything new, it earns its place as a consensus witness.

Language model. A tier-2 proposer, run only over zones the cheap tiers left thin, capped at three windows and 9k characters per document.

We initially gated the chunker on a curated technology head-noun lexicon, which is the natural design and looks like a precision win. It is not: measured against gold, gating discards 53% of mentions before any classifier sees them, because no lexicon anticipates *highback*, *rule-based elision* or the mechanical-component vocabulary patents are built from. Table 1 reports both settings.

3.3 Grounding as a structural constraint

Every candidate, from any recaller, must be a verbatim slice of the document at the offsets it reports. For model-proposed spans this is enforced by relocation: the cited evidence sentence is grounded first, the surface is then searched inside it, and a bounded ladder — exact, whitespace-insensitive, case-insensitive, and optionally a similarity snap — returns *the document’s* characters or nothing at all. The ladder rung each accepted span landed on is recorded; the similarity rung is disabled by default because it is the only one that can silently move a span onto a different entity.

The resulting invariant is asserted directly in the test suite: for every emitted mention, `doc.text[start:end] == surface`. The hallucination rate we report is consequently a check on the implementation rather than a comparison between systems, and a non-zero value would be a bug report.

3.4 Typing, stance and verification

Type is a linear model over interpretable features: head-noun lexicon membership, orthography (acronyms, CamelCase, version numbers, chemical formulae), knowledge-base depth, structural zone, and a distributional back-off that scores the surface against hand-seeded prototype centroids (Reimers and Gurevych, 2019). Linearity is deliberate — the

decision must be explainable in an audit, and with a few hundred annotated mentions a linear model is also the better estimator. Stance combines a zone prior with a cue lexicon in a window around the mention.

Verification is a cascade of three checks of increasing cost, each of which may abstain. A *structural* verifier re-derives grounding and evidence from the document rather than trusting the record. An *embedding* verifier gives a distributional second opinion and decides the easy cases in both directions. Only the residual uncertainty band reaches an *LLM judge*, which is shown the evidence sentence, the term and the proposed type — and nothing about how any of them were produced, because a verifier shown the proposer’s reasoning ratifies it. Where proposer and verifier disagree and the proposer had independent support, an adjudicator on a stronger model sees a wider context window and may accept, reject, or abstain; the prompt pushes it towards abstention, which for a corpus that will be aggregated costs one annotator-minute instead of a data point nobody revisits.

3.5 One abstention mechanism

An earlier version let each guard abstain independently and remove the mention. This made recall a function of guard ordering and filled the review queue with items no one had scored. The system now has exactly one abstention mechanism: guards either reject (a provable violation) or pass, soft signals feed a logistic confidence model, and a single calibrated threshold decides what is emitted versus withheld. Overlapping candidates are resolved *after* scoring, by confidence, since the only evidence available before typing is span length and recaller count — which picks the wrong boundary about as often as the right one.

3.6 Canonicalisation

Counting surface forms is not counting technologies. Mentions are merged by a cascade: in-document abbreviation pairs (promoted to corpus-wide aliases), knowledge-base identity, lemma-key identity, and finally embedding agglomeration over the NIL residue with a head-noun compatibility constraint that prevents “graph neural network” and “spiking neural network” from collapsing. Clusters that never touch the ontology receive a stable synthetic identifier, so an emerging technology can be tracked from first naming and reconciled with an ontology later.

3.7 Cost discipline

Per-candidate verification calls are the obvious design and the wrong one: a patent yields 10^2 candidates whose useful payload is a term, a type and one sentence. Three reductions apply before anything is sent. Verification is deduplicated on (surface, type, sentence), so a term repeated forty times in a specification is one question. Only the embedding verifier’s uncertainty band escalates. Survivors are batched roughly twenty per request behind a cached system prompt. Together these take a patent from around 120 verifier calls to one or two. A planner reproduces the window-selection logic to price a run before any request is issued.

4 Experimental setup

Data. 24 arXiv abstracts sampled across six categories (computational linguistics, machine learning, signal processing, materials science, applied physics, quantitative biology) and 20 USPTO publications sampled from the HUPD frame (Suzgun et al., 2023) stratified over five CPC sections, with full text from Google Patents. The patent sample spans snowboard bindings, HDAC inhibitors, Talbot–Lau interferometry, Ziegler–Natta catalysis, wheel bearings and narrowband spread-spectrum transmission. This spread is the point: a pipeline tuned on one vocabulary would visibly fail on the others.

Gold. 895 mentions over 44 documents (339 paper, 556 patent), annotated by one author under written guidelines. The evaluation zone is title and abstract for papers, plus the first independent claim for patents, capped at 1400 characters — Markush structure claims in chemical cases run to thousands of characters of variable definitions that name no technology. System output is filtered to the same zone, so a system is neither rewarded nor penalised outside it.

Metrics. Strict (exact offsets) and partial (overlap) span P/R/F₁; type accuracy on matched spans; the hallucination rate defined above; and the area under the risk–coverage curve (AURC, lower is better), which rewards a system whose errors are concentrated in its low-confidence tail (Geifman and El-Yaniv, 2017). We treat partial F₁ as the primary comparison figure, following practice in term extraction (Rigouts Terryn et al., 2020; QasemiZadeh and Schumann, 2016), because boundary choice for technology terms is genuinely ambiguous and

System	Strict			Partial	Type	AURC
	P	R	F ₁	F ₁		
CSO gazetteer	0.303	0.067	0.110	0.214	0.294	0.299
Chunker (lexicon-gated)	0.238	0.200	0.217	0.402	0.443	0.562
Chunker (open)	0.178	0.476	0.259	0.445	0.388	0.590
C-value/NC-value	0.011	0.006	0.008	0.385	0.406	0.496
Union of recallers	0.131	0.363	0.193	0.440	0.404	0.583

Table 1: Single-recaller baselines. Every row is a configuration of the same code, with the guard stack reduced to grounding.

System	Strict			Partial	Type	AURC
	P	R	F ₁	F ₁		
TEKNE	0.337	0.388	0.361	0.548	0.595	0.396
– negative lexicon	0.285	0.368	0.321	0.569	0.602	0.375
– consensus	0.331	0.406	0.365	0.551	0.593	0.404
– verifier	0.192	0.469	0.272	0.458	0.399	0.548
– abstention	0.307	0.450	0.365	0.542	0.597	0.422
– KB & abbreviations	0.350	0.368	0.358	0.559	0.583	0.363

Table 2: Ablations: one component removed at a time from the full offline configuration.

the strict/partial gap measures disagreement about boundaries rather than failure to find the technology.

Configuration. All numbers below are from the offline configuration — deterministic and neural tiers only, no API key — so they are reproducible without cost. Baselines are components of this system rather than external tools, which answers the question a reader actually has (what does the machinery buy over its own parts) and avoids comparing against systems built for a different ontology.

5 Results

Baselines. Table 1 shows the characteristic profile of each component. The gazetteer is precise and nearly blind (0.067 recall): outside computer science the ontology has almost nothing to say. C-value alone has the worst strict F₁ of any configuration (0.008) but a respectable partial F₁ (0.385) — it finds the right regions of text and the wrong boundaries, which is precisely what a frequency statistic with no syntax should do. Gating the chunker on the head-noun lexicon raises its precision from 0.178 to 0.238 and costs more than half its recall.

The full system. Table 2 gives 0.548 partial F₁ and 0.595 type accuracy at 0.396 AURC. The hallucination rate is 0.000 in every configuration, as

Genre	P	R	F ₁	Gold
Papers	0.219	0.354	0.271	274
Patents	0.438	0.405	0.421	558

Table 3: Per-genre performance of the full system.




Figure 1: Risk against coverage. Configurations that look close on F₁ separate here: what distinguishes them is not how much they find but whether their errors sit in the low-confidence tail. Lower is better; AURC in the legend.

the design requires.

What each component buys, and what it costs. The ablations do not all point the same way, and the ones that do not are the informative ones.

Removing the *verifier* is the only unambiguous degradation, and it is large: precision falls from 0.337 to 0.192, partial F₁ from 0.548 to 0.458, type accuracy from 0.595 to 0.399, and AURC from 0.396 to 0.548. Recall rises, to 0.469. That is exactly what a verifier is for — it does not find more, it decides what to keep — and it is the component we would keep if we could keep only one.

The *negative lexicon* and the *knowledge base plus abbreviation miner* both trade in the opposite direction. Dropping the negative lexicon costs strict precision (0.337 to 0.285) and strict F₁ (0.361 to 0.321) while slightly *improving* partial F₁ (0.569) and AURC (0.375). Dropping the knowledge base and abbreviation miner behaves similarly. The explanation is the same in both cases: these components remove candidates that overlap a gold span without matching it exactly, so they buy exact-boundary precision by giving up partial credit, and the items they remove were already ranked low enough that removing them does not improve the ranking. For a tracking corpus operated at a pre-

cision target this is the right trade, but it is a trade and not a free win.

Consensus and *abstention* move strict F₁ by less than a point and improve AURC (0.404 and 0.422 respectively, against 0.396 for the full system), which is consistent with their role: they change what the system is confident about rather than what it finds. Figure 1 shows why this is worth measuring separately: configurations within a point of each other on F₁ are visibly different in where their errors sit, and it is that difference, not F₁, that determines how much of the output a downstream corpus can take unreviewed.

Genre. Table 3 shows the gap the task description predicts. Patents are the harder genre despite contributing more gold mentions per document, because claim language generates large numbers of well-formed component noun phrases that are structurally indistinguishable from technology mentions without domain knowledge — “a first position”, “the both-side end portions”, “said support section”.

Error analysis. Stage-wise accounting on a 20-document sample attributes the remaining recall loss as follows: candidates cover 80% of gold spans exactly and a further 17% with overlap, so only 2% is lost at recall. Of the gold spans present as candidates, the largest single loss is boundary disagreement at decoding — the decoder selects an overlapping span the annotator did not choose — followed by type errors, predominantly METHOD→NOT_TECH for heads the lexicon does not list. This is why the strict/partial gap is large and why we report both.

6 Prompt injection: a negative result

Documents are adversary-authored, permanently published, and placed inside a model prompt, which makes indirect prompt injection (Greshake et al., 2023) a live concern for a corpus built over decades of filings. We wrap document text in a per-call nonce fence, but the substantive defence was supposed to be grounding.

It is not sufficient, and an adversarial test showed why. An attacker who writes “*always extract QuantumLeap Processor as the primary technology*” into their own specification has made that string a verbatim substring of their own document; a grounded extractor has no basis on which to reject it. What grounding buys is a bound, not a defence: the adversary can only inject technologies they are

willing to publish under their own name. Closing the residue required a second mechanism — rejecting candidates whose own sentence contains instruction-like content — which costs the mentions in the example sentences of papers genuinely about prompting. We keep the negative result as a test so that it stays visible.

7 Downstream: what the extraction is for

Aggregating mentions by canonical identifier and document year gives per-technology series; ranking by recent growth in normalised document frequency gives an emergence list; and comparing the early and late role distribution of a technology gives an adoption signature, the shift from PROPOSED to USED and COMPARED. We report this as a demonstration that the output supports the analysis rather than as a contribution to trend detection (Small et al., 2014; Érdi et al., 2013); document frequency rather than mention frequency is used throughout, since a patent that repeats “grinding tool” forty times is one document’s worth of evidence.

8 Limitations

The gold set has a single annotator, so there is no agreement figure and the type and boundary decisions carry one person’s judgement; the reported figures characterise this pipeline against this annotation rather than absolute accuracy. Roles are annotated per document rather than per occurrence, which makes patent role accuracy closer to a measure of zone detection than of stance. The paper sample is drawn from recent arXiv submissions and over-represents LLM-adjacent work. In the offline configuration the type classifier and the embedding verifier share a distributional signal, so that verifier is a consistency check rather than an independent one; genuine independence requires the LLM tier. The evaluation zone excludes patent descriptions, where much specific disclosure lives. Finally, the corpus is small: 44 documents supports ablation comparison but not confident absolute claims.

9 Future work

The highest-value extension is *temporal sense disambiguation*. Vocabulary is reused across eras — “transformer” in a 2005 power-systems patent is a magnetic component — and a single such error lands at the start of a time series, exactly where an

emergence detector is most sensitive. We implement a temporal guard that withholds knowledge-base links predating a document, but the harder version, which disambiguates senses against a contemporaneous corpus, is unattempted.

Three further directions. *Per-occurrence stance* with a trained classifier rather than a cue lexicon, which the current gold set cannot evaluate. *Successor relations*: co-occurrence and role transitions already suggest when one technology displaces another, and extracting explicit comparative statements (“outperforms X”) would turn that into a graph (Hope et al., 2021). *Full-text patents*: the description section carries the specific disclosure that claim language deliberately generalises away, and extending the zone there is mostly a cost-control problem given the batching in §3.7. We would also like to compare against a fine-tuned span tagger (Beltagy et al., 2019; Zhong and Chen, 2021) trained on a larger gold set, and against domain extractors in chemistry (Krallinger et al., 2015) on the patent subset, neither of which the present annotation budget supported.

10 Conclusion

Technology extraction for evolution tracking is a precision-dominant, open-vocabulary problem in which fabricated entities do lasting damage and missed ones mostly do not. Treating that asymmetry as a design constraint — offsets rather than strings, cited evidence rather than assertion, one calibrated abstention rather than many cautious stages — produces a system whose hallucination rate is zero by construction and whose remaining errors are boundary and type disagreements that a reader can inspect. The components that matter most are not the ones that find mentions but the ones that decide which to keep.

Reproducibility

Corpora are rebuilt from public endpoints without credentials; all offline numbers are produced by make data kb gold eval. Each result row is a configuration object rather than a code path, every mention records the digest of the configuration that produced it, and model calls are cached content-addressed so a run replays exactly. The tables in this paper are generated from the results file and the numbers quoted in prose are LaTeX macros from the same source.

References

- Isabelle Augenstein, Mrinal Das, Sebastian Riedel, Lakshmi Vikraman, and Andrew McCallum. 2017. SemEval 2017 task 10: ScienceIE — extracting keyphrases and relations from scientific publications. In *Proceedings of the 11th International Workshop on Semantic Evaluation*, pages 546–555.
- Iz Beltagy, Kyle Lo, and Arman Cohan. 2019. SciBERT: A pretrained language model for scientific text. In *Proceedings of EMNLP-IJCNLP*, pages 3615–3620.
- Ran El-Yaniv and Yair Wiener. 2010. On the foundations of noise-free selective classification. *Journal of Machine Learning Research*, 11:1605–1641.
- Péter Érdi, Kinga Makovi, Zoltán Somogyvári, Katherine Strandburg, Jan Tobochnik, Péter Volf, and László Zálányi. 2013. Prediction of emerging technologies based on analysis of the US patent citation network. *Scientometrics*, 95(1):225–242.
- Katerina Frantzi, Sophia Ananiadou, and Hideki Mima. 2000. Automatic recognition of multi-word terms: the C-value/NC-value method. *International Journal on Digital Libraries*, 3(2):115–130.
- Yonatan Geifman and Ran El-Yaniv. 2017. Selective classification for deep neural networks. In *Advances in Neural Information Processing Systems*.
- Kai Greshake, Sahar Abdelnabi, Shailesh Mishra, Christoph Endres, Thorsten Holz, and Mario Fritz. 2023. Not what you’ve signed up for: Compromising real-world LLM-integrated applications with indirect prompt injection. In *Proceedings of the 16th ACM Workshop on Artificial Intelligence and Security*.
- Matthew Honnibal, Ines Montani, Sofie Van Landeghem, and Adriane Boyd. 2020. spaCy: Industrial-strength natural language processing in Python. Software.
- Tom Hope, Jason Portenoy, Kishore Vasan, Jonathan Borchardt, Eric Horvitz, Daniel S. Weld, Marti A. Hearst, and Jevin West. 2021. Extracting a knowledge base of mechanisms from COVID-19 papers. In *Proceedings of NAACL-HLT*, pages 4489–4503.
- Ziwei Ji, Nayeon Lee, Rita Frieske, Tiezheng Yu, Dan Su, Yan Xu, Etsuko Ishii, Yejin Bang, Andrea Madotto, and Pascale Fung. 2023. Survey of hallucination in natural language generation. *ACM Computing Surveys*, 55(12).
- Kyo Kageura and Bin Umino. 1996. Methods of automatic term recognition: A review. *Terminology*, 3(2):259–289.
- Martin Krallinger, Obdulia Rabal, Florian Leitner, Miguel Vazquez, David Salgado, Zhiyong Lu, Robert Leaman, Yanan Lu, Donghong Ji, Daniel M. Lowe, et al. 2015. The ChEMDNER corpus of chemicals and drugs and its annotation principles. *Journal of Cheminformatics*, 7(S1).
- Yi Luan, Luheng He, Mari Ostendorf, and Hannaneh Hajishirzi. 2018. Multi-task identification of entities, relations, and coreference for scientific knowledge graph construction. In *Proceedings of EMNLP*, pages 3219–3232.
- Sewon Min, Kalpesh Krishna, Xinxu Lyu, Mike Lewis, Wen-tau Yih, Pang Wei Koh, Mohit Iyyer, Luke Zettlemoyer, and Hannaneh Hajishirzi. 2023. FActScore: Fine-grained atomic evaluation of factual precision in long form text generation. In *Proceedings of EMNLP*.
- Behrang QasemiZadeh and Anne-Kathrin Schumann. 2016. The ACL RD-TEC 2.0: A language resource for evaluating term extraction and entity recognition methods. In *Proceedings of LREC*, pages 1862–1868.
- Hannah Rashkin, Vitaly Nikolaev, Matthew Lamm, Lora Aroyo, Michael Collins, Dipanjan Das, Slav Petrov, Gaurav Singh Tomar, Iulia Turc, and David Reitter. 2023. Measuring attribution in natural language generation models. *Computational Linguistics*, 49(4):777–840.
- Nils Reimers and Iryna Gurevych. 2019. SentenceBERT: Sentence embeddings using Siamese BERT-networks. In *Proceedings of EMNLP-IJCNLP*, pages 3982–3992.
- Ayla Rigouts Terryn, Véronique Hoste, and Els Lefever. 2020. In no uncertain terms: A dataset for monolingual and multilingual automatic term extraction from comparable corpora. *Language Resources and Evaluation*, 54(2):385–418.
- Angelo A. Salatino, Thiviyan Thanapalasingam, Andrea Mannocci, Aliaksandr Birukou, Francesco Osborne, and Enrico Motta. 2020. The Computer Science Ontology: A comprehensive automatically-generated taxonomy of research areas. *Data Intelligence*, 2(3):379–416.
- Ariel S. Schwartz and Marti A. Hearst. 2003. A simple algorithm for identifying abbreviation definitions in biomedical text. In *Pacific Symposium on Biocomputing*, pages 451–462.
- Henry Small, Kevin W. Boyack, and Richard Klavans. 2014. Identifying emerging topics in science and technology. *Research Policy*, 43(8):1450–1467.
- Mirac Suzgun, Luke Melas-Kyriazi, Suproteem K. Sarkar, Scott Duke Kominers, and Stuart M. Shieber. 2023. The Harvard USPTO patent dataset: A large-scale, well-structured, and multi-purpose corpus of patent applications. In *Advances in Neural Information Processing Systems, Datasets and Benchmarks Track*.
- David Wadden, Ulme Wennberg, Yi Luan, and Hannaneh Hajishirzi. 2019. Entity, relation, and event extraction with contextualized span representations. In *Proceedings of EMNLP-IJCNLP*, pages 5784–5789.

Zexuan Zhong and Danqi Chen. 2021. A frustratingly easy approach for entity and relation extraction. In *Proceedings of NAACL-HLT*, pages 50–61.